Enhancements for Monte Carlo Tree Search in The Mario AI Framework

Emil Juul Jacobsen   Rasmus Greve
ejuu@itu.dk   ragr@itu.dk

May 22, 2013

Abstract

(Bare copy/paste-ish fra projektbasen, skal skrives om!) In this experiment we explore different implementations and enhancements of the Monte Carlo Tree Search algorithm for an AI, in order to evaluate their performance and results in the Super Mario AI Benchmark tool. We have implemented the basic MCTS algorithm in the Mario AI Framework and characterised the performance and identification of the strengths and weaknesses of the algorithm relative to the framework. We have identified a set of refinements and alterations of the algorithm and through implementation and evaluation of these individually we came up with compositions that greatly increase the performance of the AI.

1   Introduction

In this experiment we explore different implementations and enhancements of the Monte Carlo Tree Search (MCTS) algorithm for an AI, in order to evaluate their performance and results in the Super Mario AI Benchmark tool [5]. The MCTS algorithm has shown great results in various classical board games but has (to our knowledge) not been tested on a real-time physics-based game like Super Mario. Like some of the games that MCTS has proven effective in, Super Mario has a large branching factor of the state space but differs in that simulating actions is quite computationally expensive. These differences make several modifications of the core algorithm interesting for our experiment because they can help build the tree in a manner that uses the simulations more effectively.

2 Background

2.1 Monte Carlo Tree Search

Monte Carlo Tree Search has shown surprisingly good results in solving problems with large branching factors. MCTS is a tree search algorithm, where evaluation of a node is done by random sampling in the decision space until an outcome can be determined. A very nice property of MCTS is that it is an *anytime* algorithm, meaning that it can be halted when a time budget is reached and give the result that looks the most promising at the given time. Furthermore it often requires no or very little domain knowledge as a naïve implementation only require knowledge of the action space and a means of simulating the outcome of an action.

Searching in MCTS is done by iteratively by building a search tree where the nodes are game states, and the edges are actions. A node is added to the tree in each iteration and recursively, based on the reward of the new node, the values of parent nodes are updated. A single iteration of the MCTS building process consists of these four steps:

1. Tree Policy (A node to be expanded is chosen)

2. Expansion (The node is expanded by simulating the associated action)

3. Default Policy (The game is simulated with random moves until terminal)

4. Backpropagation (The result propagates up through the tree)

Modifications to this core algorithm will change one or more of these steps. We will note for each modification what steps they modify.

2.1.1 Tree Policy

First the algorithm must search through the tree and determine what existing node to expand. This is done by traversing the tree until we find a node that is not fully expanded. Modifications to this part of the algorithm generally differs in the choice between existing child nodes when looking at a fully expanded node.

2.1.2 Expansion

From the node that is not fully expanded a new child node is created corresponding to performing a specific action to the state of the current node. Differences in policies here are due to selecting the new child and how the action is performed (e.g. section 3.4).

2.1.3 Default Policy

The default policy determines what is done leading up to the final calculation of the reward for the new node. For the basic MCTS this is doing random actions until we reach terminal which could either be a win or a loss.

2.1.4 Backpropagation

The final reward calculated after performing the default policy is saved in the node, and propagated up the tree to all ancestors. This value will then be used by the tree policy in later iterations.

2.2 Upper Confidence Bounds for Trees

Choosing which node to expand can be done in multiple ways. Upper Confidence Bound (UCB) is a bandit based approach to choosing the most urgent node to expand. The quality with UCB is that it allows for prioritizing between exploration of less tried nodes, and exploitation of seemingly promising nodes.

$$UCB_j = \overline{X}_j + \sqrt{\frac{2 * \log(n)}{n_j}} \tag{1}$$

Here X_j is the average (i.e. expected) reward. n is the total number of plays and n_j is the number of times action j was tried. Since the exploration term of the equation depends on how explored the node is compared to the parent, the confidence in a node will increase steadily until the node is eventually explored. The exploitation term will, however, make sure that good nodes will be explored more frequently than less promising ones.

2.3 The Mario AI Framework

The Super Mario implementation we use to test our enhancements in, is a framework

2.3.1 The Game API

2.3.2 Performance measurement

2.4 Foundation

Vi tager simulationen fra Baumgarten

3 Approach and Enhancements

In this section we will go over the different implementations and enhancements that we have experimented with. For enhancements including a value that can be tweaked we have tried finding the optimal setting, and in the end we have combined enhancements which complement each other well.

3.1 Monte Carlo Tree Search with UCB

3.1.1 Description

To have a baseline for comparison and an implementation to build on, we started by implementing the basic Monte Carlo Tree Search algorithm with Upper Confidence Bounds for selection of nodes. Source [1]

3.1.2 Implementation

3.1.3 Behaviour

3.2 Domain knowledge

While MCTS on it's own does show some intelligent behaviour, we felt like we had to introduce some domain knowledge to the algorithm to make it perform better. Otherwise most of the enhancements to make wouldn't really show much difference, as it in any case would lead to quite a large amount of random behaviour.

3.2.1 Description

Down button   The absolute first piece of domain knowledge that we introduced was to remove the possibility of pressing the down arrow making Mario duck. In a very few cases it indeed does make sense to duck - under a bullet or a flying monster - but while it can be of use, it also doubles the action space from 16 to 32 possible actions which reduces the reachable depth significantly. Another important factor that made us disable the down button was that while the down button is pressed, pressing left or right does not make Mario move.

Specific actions   Blablabla om specifikke actions

Hole Detection   Blablabla om hole detection

3.2.2 Implementation

Down button and specific actions   Blablabla om specifikke actions

Hole detection   Blablabla om hole detection

3.2.3 Behaviour

Down button and specific actions   Blablabla om specifikke actions

Hole detection   Blablabla om hole detection

3.3   Softmax Backup

3.3.1   Description

$$exploitation = Q * maxReward + (1 - Q) * averageReward \qquad (2)$$

We use equation 2 to calculate the exploration part for the confidence of nodes

3.3.2   Implementation

3.3.3   Behaviour

3.4   Macro actions

[3]

3.4.1   Description

3.4.2   Implementation

3.4.3   Behaviour

3.5   Heuristic Partial Tree Expansion Policy

3.5.1   Description

3.5.2   Implementation

3.5.3   Behaviour

3.6   Checkpoints

3.6.1   Description

3.6.2   Implementation

3.6.3   Behaviour

3.7   (Combination)

3.7.1   Description

3.7.2   Implementation

3.7.3   Behaviour

4   Results

Method	Score
Softmax backup q = 0 (UCT)	34,162
Softmax backup q = $\frac{1}{8}$	-
Softmax backup q = $\frac{1}{4}$	34,387
Softmax backup q = $\frac{1}{2}$	34,147
Softmax backup q = 1 (Max)	26,842

Table 1: Results of using Softmax backup with different q values

Method	Avg. number of nodes	Score
MCTS w/ UCT, limit = 0	-	-
MCTS w/ UCT, limit = 1	-	-
MCTS w/ UCT, limit = 2	-	-
MCTS w/ UCT, limit = 4	-	-
MCTS w/ UCT, limit = 8	-	-
MCTS w/ UCT, limit = 16	-	-
MCTS w/ UCT, limit = ∞	-	-

Table 2: Results of using UCT with a different limit for random moves

Her er noget mere tekst, reference til figur 1

Method	Score
MCTS w/ UCT, cp = 0	-
MCTS w/ UCT, cp = $\frac{1.5}{8}$	-
MCTS w/ UCT, cp = $\frac{1}{4}$	-
MCTS w/ UCT, cp = $\frac{1}{3}$	-
MCTS w/ UCT, cp = $\frac{1}{2}$	-
MCTS w/ UCT, cp = $\frac{1}{\sqrt{2}}$	-
MCTS w/ UCT, cp = 2	-
MCTS w/ UCT, cp = 5	-
MCTS w/ UCT, cp = 10	-

Table 3: Results of using UCT with different weight between exploration and exploitation

Figure 1: Mario being followed

Method	Score
UCT (Softmax q = 0)	-
Max (Softmax q = 1)	-
Softmax	-
Limited Actions	-
Hole Detection	-
Macro Actions	-
H. Partial Expansion	-
Checkpoints	-
(Combination)	-

Table 4: Results from all the different enhancements

5   Conclusion

References

[1] C. B. Browne, E. Powley, D. Whitehouse, S. M. Lucas, P. I. Cowling, P. Rohlfshagen, S. Tavener, D. Perez, S. Samothrakis, and S. Colton, "A survey of monte carlo tree search methods," Computational Intelligence and AI in Games, IEEE Transactions on, vol. 4, no. 1, pp. 1–43, 2012.

[2] S. R. K. Branavan, D. Silver, and R. Barzilay, "Non-linear monte-carlo search in civilization ii," in Proceedings of the Twenty-Second international joint conference on Artificial Intelligence-Volume Volume Three, 2011, pp. 2404–2410.

[3] E. J. Powley, D. Whitehouse, and P. I. Cowling, "Monte Carlo Tree Search with macro-actions and heuristic route planning for the Physical Travelling Salesman Problem," in Computational Intelligence and Games (CIG), 2012 IEEE Conference on, 2012, pp. 234–241.

[4] T. Pepels and M. H. Winands, "Enhancements for Monte-Carlo Tree Search in Ms Pac-Man," in Computational Intelligence and Games (CIG), 2012 IEEE Conference on, 2012, pp. 265–272.

[5] S. Karakovskiy and J. Togelius, "The mario ai benchmark and competitions," Computational Intelligence and AI in Games, IEEE Transactions on, vol. 4, no. 1, pp. 55–67, 2012.